

conv_filter_visualization_v2

April 18, 2018

1 What CNN Sees

This notebook is based solely from the content presented in <https://blog.keras.io/how-convolutional-neural-networks-see-the-world.html> by *Francois Chollet* and https://github.com/keras-team/keras/blob/master/examples/conv_filter_visualization.py.

Some extra content have been added from the original content to aid my understanding of using **Keras**.

This notebook visualises what deep convolutional neural networks see from the images we feed them.

We will start by defining the VGG16 model in Keras by

- (1) Import Keras applications using the command `from keras import applications` Keras Applications are deep learning models that are made available alongside pre-trained weights.
- (2) Instantiate a VGG16 model (performing this, weights are downloaded automatically to `~/.keras/models/`).

```
keras.applications.vgg16.VGG16(include_top=True, weights='imagenet', input_tensor=None, inp
```

- `include_top`: whether to include the 3 fully-connected layers at the top (end) of the network. In this code, we set `include_top = False`. This is because adding the FC layers requires us to use a fixed input size for the model (i.e. 224x244, the original ImageNet format).
- `weights`: it can be `None` for random initialisation or `imagenet` for pre-trained weights on ImageNet

```
In [2]: from keras import applications
```

```
# Build the VGG16 network
model = applications.VGG16(include_top = False,
                           weights = 'imagenet')
model.summary()
```

```
-----
Layer (type)                   Output Shape          Param #
=====
input_2 (InputLayer)          (None, None, None, 3)    0
```

```

-----
block1_conv1 (Conv2D)      (None, None, None, 64)    1792
-----
block1_conv2 (Conv2D)      (None, None, None, 64)    36928
-----
block1_pool (MaxPooling2D) (None, None, None, 64)    0
-----
block2_conv1 (Conv2D)      (None, None, None, 128)   73856
-----
block2_conv2 (Conv2D)      (None, None, None, 128)   147584
-----
block2_pool (MaxPooling2D) (None, None, None, 128)   0
-----
block3_conv1 (Conv2D)      (None, None, None, 256)   295168
-----
block3_conv2 (Conv2D)      (None, None, None, 256)   590080
-----
block3_conv3 (Conv2D)      (None, None, None, 256)   590080
-----
block3_pool (MaxPooling2D) (None, None, None, 256)   0
-----
block4_conv1 (Conv2D)      (None, None, None, 512)   1180160
-----
block4_conv2 (Conv2D)      (None, None, None, 512)   2359808
-----
block4_conv3 (Conv2D)      (None, None, None, 512)   2359808
-----
block4_pool (MaxPooling2D) (None, None, None, 512)   0
-----
block5_conv1 (Conv2D)      (None, None, None, 512)   2359808
-----
block5_conv2 (Conv2D)      (None, None, None, 512)   2359808
-----
block5_conv3 (Conv2D)      (None, None, None, 512)   2359808
-----
block5_pool (MaxPooling2D) (None, None, None, 512)   0
=====
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0
-----

```

Create a dictionary for (layer.name, layer) using layer.name as key because each layer has a unique name. We could list the layer name by

```

for layer in model.layers:
    print(layer.name)

```

```
#print('input_shape: {}\\toutput_shape: {}'.format(layer.input_shape,
                                                    layer.output_shape))
```

This results in a list of all layer in VGG16 network (notice the missing FC layers as the results of setting `include_top = False`).

```
input_1 block1_conv1 block1_conv2 block1_pool block2_conv1 block2_conv2 block2_pool
block3_conv1 block3_conv2 block3_conv3 block3_pool block4_conv1 block4_conv2 block4_conv3
block4_pool block5_conv1 block5_conv2 block5_conv3 block5_pool
```

```
In [3]: layer_dict = dict([(layer.name, layer) for layer in model.layers])
```

Next, define a **loss function** that will seek to **maximize the activation** of a specific filter (`filter_index`) in a specific layer (`layer_name`). We do this via a Keras backend function, which allows our code to run both on top of TensorFlow and Theano.

```
from keras import backend as K
```

`layer_name` can be set as `layer.name` of any layer with filter. For example,

```
layer_name = 'block5_conv3'
```

`filter_index` can be set to any integer that corresponds to the filter index in that layer. For example, there are 512 filters for `block5_conv3`. Thus, we can set the index to any value between 0 and 511.

Note that here we are seeking a **maximum value**. That is we are performing **gradient ascent**. This is the opposite process of **gradient descent** in which the latter seeks to minimise the value (typically loss) while *gradient ascent* seeks to maximise the value (for example, profit).

Recap * Gradient descent is used to find a minimum of a function. At a point, we take **negative** steps proportional to the gradient. Usually, it is $-\alpha * \text{gradient}$, where α is learning rate. * Gradient ascent is used to find a maximum of a function. We take **positive** steps proportional to the gradient. Usually, it is $\alpha * \text{gradient}$.

```
In [4]: from keras import backend as K
```

```
layer_name = 'block5_conv3'
filter_index = 0 # can be any integer from 0 to 511, as there are 512 filters
```

```
# build a loss function that maximises the activation of the nth filter of the chosen layer
# 1. obtain the output of the selected layer 'layer_name'
```

```
layer_output = layer_dict[layer_name].output
# 2. Get the mean of the tensor at filter_index
if K.image_data_format() == 'channels_first':
    loss = K.mean(layer_output[:, filter_index, :, :])
else:
    loss = K.mean(layer_output[:, :, :, filter_index])
```

```
# Define the placeholder for the input images
input_img = model.input
```

```

# compute the gradient of the input picture wrt this loss
grads = K.gradients(loss, input_img)[0]

# normalise the gradient
small_value = 1e-5 # or K.epsilon()
grads /= (K.sqrt(K.mean(K.square(grads))) + small_value)

# This function returns the loss and grads given the input picture
iterate = K.function([input_img], [loss, grads])

```

In the code above, **the gradient of the pixels of the input image is normalised**. This avoids very small and very large gradients which ensure smooth gradient ascent process. Next, we use the Keras function we defined to do **gradient ascent** in the input space, with regard to our filter activation loss. As mentioned in the previous cell, *gradient ascent is used to find a maximum of a function*, although it can be local maximum. We take **positive steps** proportional to the gradient. For the computation, that is $> \alpha \text{gradient}$ where α is learning rate.

```

In [6]: import numpy as np
        from scipy.misc import imsave

        # Let's define dimensions of the generated pictures for each filter.
        img_width = 128
        img_height = 128

        # Let's start from a gray image with some noise
        if K.image_data_format() == 'channels_first':
            input_img_data = np.random.random((1, 3, img_width, img_height))
        else:
            input_img_data = np.random.random((1, img_width, img_height, 3))
        input_img_data = (input_img_data - 0.5) * 20 + 128

        imsave('gray_image.png', input_img_data[0,:,:,:])
        print(input_img_data.shape)

        # run gradient ascent for N steps
        N = 20
        step = 1. # step size for gradient ascent
        for i in range(N):
            loss_value, grads_value = iterate([input_img_data])
            input_img_data += grads_value * step

            print('Current loss value:', loss_value)
            if loss_value <= 0.:
                # some filter get stuck to 0, skip them
                break

```

```
(1, 128, 128, 3)
```

Current loss value: 0.0

```
/home/supanee/tensorflow/lib/python3.5/site-packages/ipykernel_launcher.py:15: DeprecationWarning: `imsave` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0. Use ``imageio.imwrite`` instead.  
from ipykernel import kernelapp as app
```

Now we can extract and display the generated input.

```
In [7]: # util function to convert a tensor into a valid image  
def deprocess_image(x):  
    # normalise tensor: center at 0., and std at 0.1  
    x -= x.mean()  
    x /= (x.std() + small_value)  
    x *= 0.1  
  
    # clip to [0, 1]  
    x += 0.5  
    x = np.clip(x, 0, 1)  
  
    # convert to RGB array  
    x *= 255  
    if K.image_data_format() == 'channels_first': # c,h,w  
        x = x.transpose((1, 2, 0)) # results in h,w,c  
    x = np.clip(x, 0, 255).astype('uint8')  
    return x
```

Time to run and get the results!

```
In [8]: img = input_img_data[0]  
img = deprocess_image(img)  
imsave('%s_filter_%d.png' % (layer_name, filter_index), img)
```

```
/home/supanee/tensorflow/lib/python3.5/site-packages/ipykernel_launcher.py:3: DeprecationWarning: `imsave` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0. Use ``imageio.imwrite`` instead.
```

This is separate from the ipykernel package so we can avoid doing imports until



Input image:

Result:

